

```
In [ ]: # Import the needed library
import pandas as pd
```

Create Canned data

Canned data is hard coded within the program

Create a key:value collection of series to use to populate the dataframe for testing

```
In [ ]: data = {'Month': pd.Series(['January', 'February', 'March', 'April', 'May', 'June', 'July',
                                   'Rainfall': pd.Series([1.65, 1.25, 1.94, 2.75, 3.14, 3.65, 5.05, 1.50, 1.33, 0.07, 0.50, 2.30])
               }
```

```
In [ ]: # creates a Pandas DataFrame with two columns 'Month' and 'Rainfall',
df = pd.DataFrame(data)
df.shape
```

```
Out[ ]: (12, 2)
```

```
In [ ]: # print("Our data frame:")
# print(df, "\n")
df
```

```
Out[ ]:
```

	Month	Rainfall
0	January	1.65
1	February	1.25
2	March	1.94
3	April	2.75
4	May	3.14
5	June	3.65
6	July	5.05
7	August	1.50
8	September	1.33
9	October	0.07
10	November	0.50
11	December	2.30

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12 entries, 0 to 11
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Month       12 non-null    object
1   Rainfall    12 non-null    float64
dtypes: float64(1), object(1)
memory usage: 320.0+ bytes
```

Pandas Series is a single dimension array

Pandas dataframe is a two-dimensional array, like a spreadsheet. Our df consists of 2 rows of series(months and rainfall)

```
In [ ]: # add 1 row to the above data, with NAN and zero values
data_1 = {'Month': pd.Series(['January', 'February', 'March', 'April', 'May', 'June', 'July', 'August', 'September', 'October', 'November', 'December']),
          'Rainfall': pd.Series([1.65, 1.25, 1.94, 2.75, 3.14, 3.65, 5.05, 1.50, 1.33, 0.07, 0.50, 2.30]),
          'Temperature': pd.Series([3, 10, 15, 20, 75, "", 30, 1, 33, "", 32, 7])
          }
```

```
In [ ]: #change to df
df = pd.DataFrame(data_1)

#SAVE df as a csv file on the same path as "Processing data with Python.ipynb"
df.to_csv('rainfall.csv', index = 0)

#read a .csv file
df_rainfall = pd.read_csv('rainfall.csv')
df_rainfall
```

```
Out[ ]:
```

	Month	Rainfall	Temperature
0	January	1.65	3
1	February	1.25	10
2	March	1.94	15
3	April	2.75	20
4	May	3.14	75
5	June	3.65	NaN
6	July	5.05	30
7	August	1.50	1
8	September	1.33	33
9	October	0.07	
10	November	0.50	32
11	December	2.30	7

```
In [ ]: #to read .json file
df_json = pd.read_json('data.json')
```

```
print("Our data frame from JSON file:")
print(df_json, "\n")
```

Our data frame from JSON file:

	Month	Rainfall	Temperature
0	January	1.650	3.0
1	February	1.250	10.0
2	March	1.940	15.0
3	April	2.750	20.0
4	May	2.750	25.0
5	June	3.645	24.0
6	July	5.500	30.0
7	August	1.000	1.0
8	September	1.300	33.0
9	October	NaN	NaN
10	November	0.500	32.0
11	December	2.300	2.3

Cleaning Data:

One of the most important tasks in processing data.

Data needs to be consistent to be reliably analyzed.

Cleaning involves parsing the data detecting 'bad' or missing data

```
In [ ]: # October is NaN value, so
# To not break the algorithm we will replace NaN with the average temperature value
# Calculate the average temperature value
average_temp = df_json['Temperature'].mean()
average_temp
```

```
Out[ ]: 17.754545454545454
```

```
In [ ]: # Replace NaN value by rounded average_temp value
df_zeros = df_json.fillna(round(average_temp,0))
print("Our data with zero values: ")
print(df_zeros)
#Zero can skew the data so we should remove invalid data
# so, we will not use this data later on
```

Our data with zero values:

	Month	Rainfall	Temperature
0	January	1.650	3.0
1	February	1.250	10.0
2	March	1.940	15.0
3	April	2.750	20.0
4	May	2.750	25.0
5	June	3.645	24.0
6	July	5.500	30.0
7	August	1.000	1.0
8	September	1.300	33.0
9	October	18.000	18.0
10	November	0.500	32.0
11	December	2.300	2.3

```
In [ ]: #remove rows with the missing values
df_cleaned = df_zeros.dropna()
print("Our data with dropped values: \n", df_cleaned)
```

Our data with dropped values:

	Month	Rainfall	Temperature
0	January	1.650	3.0
1	February	1.250	10.0
2	March	1.940	15.0
3	April	2.750	20.0
4	May	2.750	25.0
5	June	3.645	24.0
6	July	5.500	30.0
7	August	1.000	1.0
8	September	1.300	33.0
9	October	18.000	18.0
10	November	0.500	32.0
11	December	2.300	2.3

```
In [ ]: #create a count of all rows containing Nans to check data before cleaning
count = 0
for index, row in df_json.iterrows():
    if any(row.isnull()):
        count = count + 1

print("\n JSON file have", str(count), "rows with Nans")
```

JSON file have 1 rows with Nans

```
In [ ]: #check the number of NAN rows in the cleaned data
count = 0
for index, row in df_cleaned.iterrows():
    if any(row.isnull()):
        count = count + 1

print("\n Number of rows with Nans: " + str(count))
```

Number of rows with Nans: 0

```
In [ ]: df_cleaned = df_clean.sort_index()
df_cleaned
```

Out[]:

	Month	Rainfall	Temperature
0	January	1.650	3.0
1	February	1.250	10.0
2	March	1.940	15.0
3	April	2.750	20.0
4	May	2.750	25.0
5	June	3.645	24.0
6	July	5.500	30.0
7	August	1.000	1.0
8	September	1.300	33.0
9	October	18.000	18.0
10	November	0.500	32.0
11	December	2.300	2.3

Statistical Analysis

Mean = the average of a set of numbers.

Median = The middle value in a sorted set of numbers.

Standard deviation = How much each value differs from the mean. Can be used to detect outliers.

Mode = The most common value in a list of data.

Pandas easily perform these functions!

In []:

```
print("Mean: ")
print(df_cleaned.mean())

print(" ")
print("Median : ")
print(df_cleaned.median())

#print(" ")
print("\n", "Standard Deviation: ")
print(df_cleaned.std())

#The mode here is wrong but will do it for explanation purposes
#rainfall has repeated value of 2.75
#temperature are all unique
print("\n", "Mode : ")
print(df_cleaned.mode())
```

```
Mean:
Rainfall      3.54875
Temperature   17.77500
dtype: float64
```

```
Median :
Rainfall      2.12
Temperature   19.00
dtype: float64
```

```
Standard Deviation:
Rainfall      4.746445
Temperature   11.626312
dtype: float64
```

```
Mode :
      Month  Rainfall  Temperature
0      April      2.75          1.0
1      August      NaN          2.3
2     December      NaN          3.0
3     February      NaN         10.0
4      January      NaN         15.0
5         July      NaN         18.0
6         June      NaN         20.0
7         March      NaN         24.0
8          May      NaN         25.0
9     November      NaN         30.0
10    October      NaN         32.0
11  September      NaN         33.0
```

```
C:\Users\ADMIN\AppData\Local\Temp\ipykernel_11364\1655336313.py:2: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
```

```
    print(df_cleaned.mean())
```

```
C:\Users\ADMIN\AppData\Local\Temp\ipykernel_11364\1655336313.py:6: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
```

```
    print(df_cleaned.median())
```

```
C:\Users\ADMIN\AppData\Local\Temp\ipykernel_11364\1655336313.py:10: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
```

```
    print(df_cleaned.std())
```

```
In [ ]: df_cleaned.describe()
```

Out[]:

	Rainfall	Temperature
count	12.000000	12.000000
mean	3.548750	17.775000
std	4.746445	11.626312
min	0.500000	1.000000
25%	1.287500	8.250000
50%	2.120000	19.000000
75%	2.973750	26.250000
max	18.000000	33.000000

Selecting Parts of a Dataframe

Indexing

Select single columns using a column name(temperature). Returns a series.

Example: `df_clean['Temperature']`

Select multiple columns using column names. Must specify a list of column names.

Example: `df_clean[['Temperature', 'Rainfall']]`

iloc and loc

Select a certain row number using `iloc`:

Example: `print("Third row \n", df_clean.iloc[2])`

Select a certain row using a certain value:

Example: `print("\n Third row \n", dfIndexed.loc['March']);`

```
In [ ]: print("Temperature column: \n", df_cleaned['Temperature'])
print("\n Temperature and Rainfall column: \n",df_cleaned[['Temperature', 'Rainfall']])

print("\n Third row \n", df_cleaned.iloc[2])
```

Temperature column:

0	3.0
1	10.0
2	15.0
3	20.0
4	25.0
5	24.0
6	30.0
7	1.0
8	33.0
9	18.0
10	32.0
11	2.3

Name: Temperature, dtype: float64

Temperature and Rainfall column:

	Temperature	Rainfall
0	3.0	1.650
1	10.0	1.250
2	15.0	1.940
3	20.0	2.750
4	25.0	2.750
5	24.0	3.645
6	30.0	5.500
7	1.0	1.000
8	33.0	1.300
9	18.0	18.000
10	32.0	0.500
11	2.3	2.300

Third row

Month	March
Rainfall	1.94
Temperature	15.0

Name: 2, dtype: object

```
In [ ]: #To use loc, we require a properly indexed dataframe
index = df_cleaned['Month']
dfIndexed = df_cleaned.set_index(index)

print("\n Third row \n", dfIndexed.loc['March'])
```

Third row

Month	March
Rainfall	1.94
Temperature	15.0

Name: March, dtype: object

```
In [ ]: #Print the rainfall and mean for the first 4 months
rainfall = df_cleaned['Rainfall'][0:4]
print(rainfall, "\n")
print("The mean of rainfall is:", round(rainfall.mean(),2))
```

0	1.65
1	1.25
2	1.94
3	2.75

Name: Rainfall, dtype: float64

The mean of rainfall is: 1.9


```
In [ ]: #Print the rainfall and mean for the first few months
print("\n Just Temperature and rainfall data ")
df_TempRain = df_cleaned[['Temperature','Rainfall']]
print(df_TempRain, "\n")
print("The mean of Temperature and Rainfall is: \n", df_TempRain.mean(), "\n")
```

Just Temperature and rainfall data

	Temperature	Rainfall
0	3.0	1.650
1	10.0	1.250
2	15.0	1.940
3	20.0	2.750
4	25.0	2.750
5	24.0	3.645
6	30.0	5.500
7	1.0	1.000
8	33.0	1.300
9	18.0	18.000
10	32.0	0.500
11	2.3	2.300

The mean of Temperature and Rainfall is:

Temperature	17.77500
Rainfall	3.54875

dtype: float64

Many thanks to MarwaEshra and Instructor (David Dalsveen)